

# Red Hat

## OCP 4.14 Stateless Ingress Node Firewall

Presenter

Mohamed S. Mahmoud

## What does it do

Secure OpenShift nodes from external (DOS) attacks by configuring specific **stateless** policies

## How does it work

- Ingress node firewall operator uses CR to deploy and configure ingress node firewall rules
- Webhook validates the configuration and avoids malformed configuration from locking up OCP cluster access (fail safe rules)
- XDP+eBPF apply ingress node firewall rules, parse packets, process rule actions, update statistics and generate syslog events for dropped packets

## Webhook - Ingress firewall Fail Safe rules

- Prevents users from locking up cluster access
  - for example by denying TCP traffic to the API server port 6443
- Webhook verification refuses **any** rules that match against critical ports for TCP and UDP with **deny** action

Source CIDR	Protocol	Port	Purpose
0.0.0.0/0	TCP	22	SSH
0.0.0.0/0	UDP	68	DHCP
0.0.0.0/0	TCP	6443	Kubernetes API server access
0.0.0.0/0	TCP	2379/2380	Etcd access
0.0.0.0/0	TCP	10250	Kubelet
0.0.0.0/0	TCP	10259	Kube-scheduler
0.0.0.0/0	TCP	10257	Kube-controller-manager

## Why XDP + eBPF

eBPF provides a very flexible mechanism to allow early detection and packet filtering

eXpress Data Path (XDP) is a feature in the Linux kernel which ...

- allows users to execute a user-supplied eBPF program when a packet is received on a network interface (NIC)
  - **flexibility**
- allows attachment of an eBPF program at the earliest networking driver stage
  - **high performance**
- allows processing of accepted packets via the Linux kernel
  - **no need to reinvent the wheel**

# How to deploy Ingress Node Firewall

CONFIDENTIAL Designator

## Enable subscription

```
cat <<EOF | oc apply -f -
kind: Namespace
metadata:
  name: openshift-ingress-node-firewall
---
apiVersion: operators.coreos.com/v1
kind: OperatorGroup
metadata:
  name: ingress-node-firewall
  namespace: openshift-ingress-node-firewall
spec:
  upgradeStrategy: Default
---
apiVersion: operators.coreos.com/v1alpha1
kind: Subscription
  name: ingress-node-firewall-sub
  namespace: openshift-ingress-node-firewall
spec:
  channel: alpha
  name: ingress-node-firewall
  source: ingress-node-firewall
  sourceNamespace: openshift-ingress-node-firewall
EOF
```

## Deploy IngressNodeFirewallConfig

```
cat <<EOF | oc apply -f -
apiVersion: ingressnodefirewall.openshift.io/v1alpha1
kind: IngressNodeFirewallConfig
metadata:
  name: ingressnodefirewallconfig
  namespace: openshift-ingress-node-firewall
spec:
  nodeSelector:
    node-role.kubernetes.io/worker: ""
EOF
```

## Check installed CRDs and running pods

### **oc get crds | grep ingressnodefirewall**

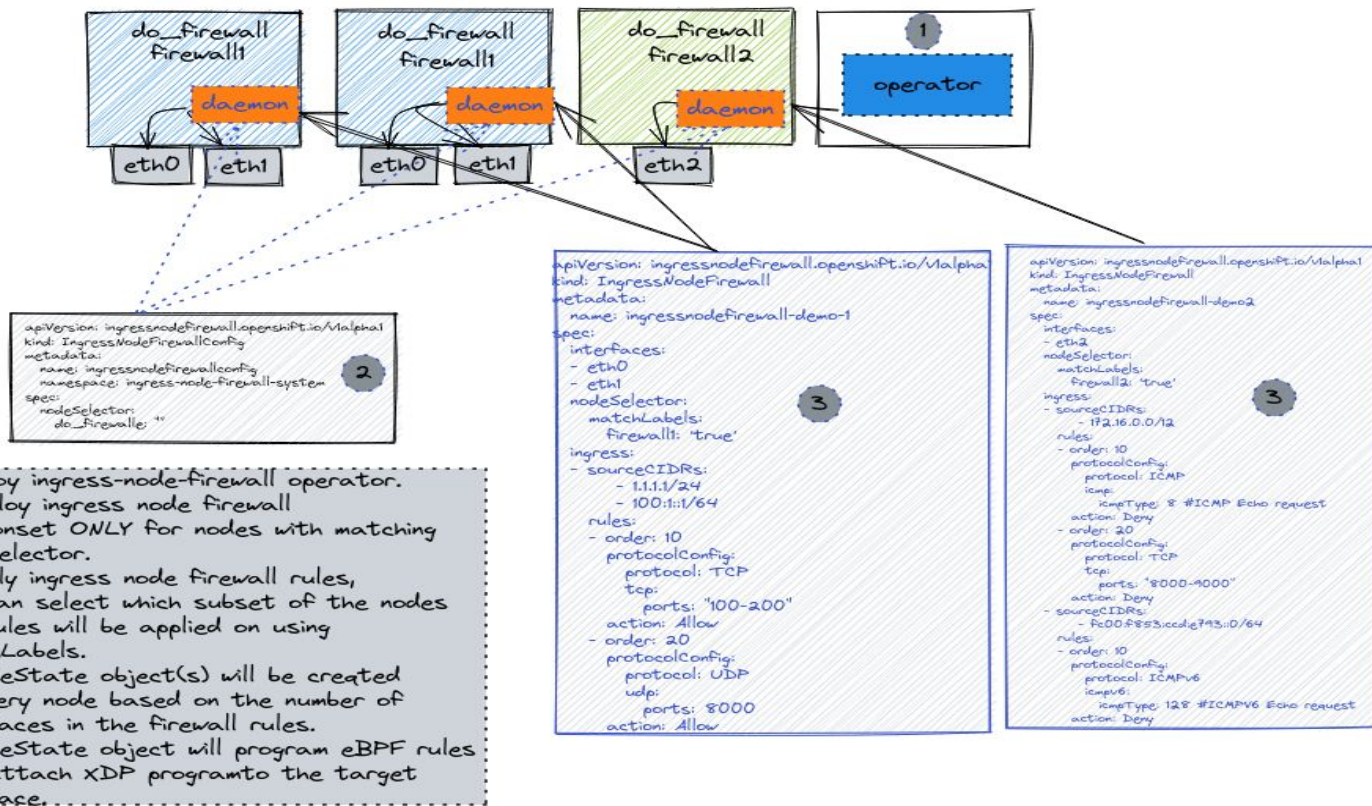
```
ingressnodefirewallconfigs.ingressnodefirewall.openshift.io    2022-08-25T10:03:01Z
ingressnodefirewallnodestates.ingressnodefirewall.openshift.io 2022-08-25T10:03:00Z
ingressnodefirewalls.ingressnodefirewall.openshift.io          2022-08-25T10:03:00Z
```

### **oc get pods -n openshift-ingress-node-firewall**

NAME	READY	STATUS	RESTARTS	AGE
ingress-node-firewall-controller-manager-656dc8dc7-cr8g5	2/2	Running	0	5d21h
ingress-node-firewall-daemon-pqx56	3/3	Running	0	5d21h

# Operator and configuration overview

CONFIDENTIAL Designator





# Ingress Firewall rules (1/5)

CONFIDENTIAL Designator

```
apiVersion: ingressnodefirewall.openshift.io/v1alpha1
kind: IngressNodeFirewall
metadata:
  name: ingressnodefirewall-demo2
spec:
  interfaces:
  - eth0
  nodeSelector:
    matchLabels:
      do-node-ingress-firewall: 'true'
  ingress:
  - sourceCIDRs:
    - 172.16.0.0/12
    rules:
    - order: 10
      protocolConfig:
        protocol: ICMP
        icmp:
          icmpType: 8 #ICMP Echo request
      action: Deny
    - order: 20
      protocolConfig:
        protocol: TCP
        tcp:
          ports: "8000-9000"
      action: Deny
```

# Ingress Firewall rules (2/5)

(...)

spec:

interfaces:

- eth0

nodeSelector:

matchLabels:

do-node-ingress-firewall: 'true'

(...)

- Specify the interfaces that rules apply to
- Using nodeSelector matchLabels user can select which nodes to apply the firewall rules to

# Ingress Firewall rules (3/5)

(...)

ingress:

```
- sourceCIDRs:  
  - 172.16.0.0/12
```

rules:

```
- order: 10
```

protocolConfig:

protocol: ICMP

icmp:

icmpType: 8 #ICMP Echo request

action: Deny

```
- order: 20
```

protocolConfig:

protocol: TCP

tcp:

ports: "8000-9000"

action: Deny

- User can configure multiple CIDRs from different address-families up to 1000 different CIDRs.
- Ingress firewall rules are ordered starting from 1 for each sourceCIDR(s) with up to 100 rules per CIDR.

# Ingress Firewall rules (4/5)

(...)

ingress:

- sourceCIDRs:
  - 172.16.0.0/12

rules:

- order: 10

protocolConfig:

protocol: ICMP

icmp:

icmpType: 8 #ICMP Echo request

action: Deny

- order: 20

protocolConfig:

protocol: TCP

tcp:

ports: "8000-9000"

action: Deny

- Supported protocols are TCP, UDP, SCTP, ICMP and ICMPv6
- For ICMP and ICMPv6 rules can match against ICMP/ICMPv6 type and/or code.
- Ingress firewall actions are either **Allow** or **Deny**
- For protocols TCP, UDP and SCTP rule can match against single DstPort or range of ports using “start-end” format

# Ingress Firewall rules (5/5)

- Matching filters are optional and ignored if they are not specified
- Ingress firewall rules get verified using verification webhook to fail any invalid configuration
- Allow and Deny packets and bytes count statistics is available per rule
- For dropped packets by XDP program an event will be emitted to syslog which includes event header (ruleId, Interface the packet came in on, packet length in bytes including L2 header) and up to 256 bytes from the packet header

# Zero Trust Ingress Firewall rules

CONFIDENTIAL Designator

```
apiVersion: ingressnodefirewall.openshift.io/v1alpha1
kind: IngressNodeFirewall
metadata:
  name: ingressnodefirewall-zero-trust
spec:
  interfaces:
  - eth1
  nodeSelector:
    matchLabels:
      do-node-ingress-firewall: 'true'
  ingress:
    - sourceCIDRs:
      - 0.0.0.0/0
      rules:
        - order: 10
          protocolConfig:
            protocol: TCP
            tcp:
              ports: 22
          action: Allow
        - order: 20
          action: Deny
```

Match any CIDR

Normally used with Multi Interfaces clusters. Where user wanted to drop all traffic on specific interface except maybe SSH for additional security.

You are logged in as a temporary administrative user. Update the [cluster OAuth configuration](#) to allow others to log in.

Project: All Projects

## OperatorHub

Discover Operators from the Kubernetes community and Red Hat partners, curated by Red Hat. You can purchase commercial software through [Red Hat Marketplace](#). You can install Operators on your clusters to provide optional add-ons to your developers. After installation, the Operator capabilities will appear in the [Developer Catalog](#) providing a self-service experience.

All Items

**Production Operators**

**AWS Load Balancer Operator**  
provided by Red Hat  
 Operator to simplify management of aws-load-balancer-controller  
 ### Prerequisites for installatio...

**Red Hat**

**AWS Load Balancer Operator**  
provided by Red Hat  
 Operator to simplify management of aws-load-balancer-controller  
 ### Prerequisites for installatio...

**Certified**

**Citrix ADC CPX Istio Sidecar Injector Operator**  
provided by citrix  
 In Istio service mesh, a sidecar proxy runs alongside application pods and it intercepts and...

**Marketplace**

**Citrix ADC CPX with Ingress Controller**  
provided by Citrix  
 Citrix ADC CPX is a container-based application delivery controller that can be provisione...

**Certified**

**Citrix ADC CPX with Ingress Controller**  
provided by Citrix  
 Citrix ADC CPX is a container-based application delivery controller that can be provisione...

**Certified**

**Citrix ADC Istio Ingress Gateway Operator**  
provided by citrix  
 An Istio ingress gateway acts as an entry point for the incoming traffic. Citrix ADC CPX, MPX, or...

**Certified**

**Citrix Ingress Controller**  
provided by Citrix  
 Citrix provides an ingress controller for Citrix ADC MPX (hardware), Citrix ADC VPX...

**Marketplace**

**Citrix Ingress Controller**  
provided by Citrix  
 Citrix provides an ingress controller for Citrix ADC MPX (hardware), Citrix ADC VPX...

**Certified**

**IBM Security Verify Operator**  
provided by IBM  
 The IBM Security Verify operator can consistently enforce policy-driven security, including...

**Production Operators**

**Ingress Node Firewall Operator**  
provided by Redhat  
 The Ingress node firewall...  
 Installed

**Certified**

**Kong**  
provided by Kong Inc.  
 Kong Operator for Kubernetes. Technical preview.

**Certified**

**Ngix Ingress Operator**  
provided by NGINX Inc.  
 The NGINX Ingress Operator is a Kubernetes/OpenShift component which deploys and...

**Source**

- ☐ Red Hat (1)
- ☐ Certified (7)
- ☐ Community (0)
- ☐ Marketplace (2)
- ☐ Production Operators (2)

**Provider**

- ☐ Red Hat (2)
- ☐ APIMatic.io (0)
- ☐ Aerospike (0)
- ☐ Alvearie (0)

AI/Machine Learning

Application Runtime

Big Data

Cloud Provider

Database

Developer Tools

Development Tools

Drivers and plugins

Integration & Delivery

Logging & Tracing

Modernization & Migration

Monitoring

Networking

OpenShift Optional

Security

Storage

Streaming & Messaging

Other

Administrator

Home

Operators

OperatorHub

Installed Operators

Workloads

Networking

Storage

Builds

Observe

Compute


User Management

Administration

You are logged in as a temporary administrative user. Update the [cluster OAuth configuration](#) to allow others to log in.

Project: openshift-ingress-node-firewall


[Installed Operators](#) > Operator details



**Ingress Node Firewall Operator**  
 4.12.0-202212061458 provided by RedHat


Actions

[Details](#)
[YAML](#)
[Subscription](#)
[Events](#)
[All instances](#)
[IngressNodeFirewallConfig](#)
[IngressNodeFirewallNodeState](#)
[IngressNodeFirewall](#)

Provided APIs

 **IngressNodeFirewallConfig**  
 Not available  
[Create instance](#)

 **IngressNodeFirewallNodeState**  
 Not available  
[Create instance](#)

 **IngressNodeFirewall**  
 Not available  
[Create instance](#)

Provider

RedHat

Created at

Dec 13, 2022, 10:09 AM

Links

Ingress Node Firewall  
<https://github.com/openshift/ingress-node-firewall>

Maintainers

msheriff234  
 mmahmoud@redhat.com

Description

The Ingress node firewall operator, manages the configuration of cluster scoped Ingress Node firewall rules for Kubernetes.

ClusterServiceVersion details

Name

ingress-node-firewall.v4.12.0-202212061458

Namespace

openshift-ingress-node-firewall

Labels

operators.coreos.com/ingress-node-firewall.openshift-ingress-node-firewall

Annotations

17 annotations

Managed Namespaces

All Namespaces

Created at

Dec 13, 2022, 10:09 AM

Status

Succeeded

Status reason

install strategy completed with no errors

Operator Deployments

ingress-node-firewall-controller-manager


Operator ServiceAccounts

ingress-node-firewall-controller-manager


OperatorGroup





openshift-ingress-node-firewall-j5kbp

16

 Red Hat



 **Red Hat**  
OpenShift

  3   kube:admin ▾

Administrator ▾

Home ▸

Operators ▾

OperatorHub

Installed Operators

Workloads ▸

Networking ▸

Storage ▸

Builds ▸

Observe ▸

Compute ▸

User Management ▸

Administration ▸


Project: openshift-ingress-node-firewall ▾


You are logged in as a temporary administrative user. Update the [cluster OAuth configuration](#) to allow others to log in.

Create IngressNodeFirewallConfig

Create by completing the form. Default values may be provided by the Operator authors.

Configure via: ☒ Form view ☐ YAML view

 Note: Some fields may not be represented in this form. Please select "YAML View" for full control of object creation.

 provided by Red Hat  
Not available

Name \*

ingressnodefirewallconfig

Labels

app=frontend

Create

Cancel

You are logged in as a temporary administrative user. Update the cluster OAuth configuration to allow others to log in.

sourceCIDs: [sourceCIDs](#)

sourceCIDRs defines the origin of packets that FirewallProtocolRules will be applied to.

Value •

fc00:f853:ccd:e793::0/64

[Add source CDRs](#)

rules

rules is a list of per protocol ingress node firewall rules.

order •

10

order defines the order of execution of ingress firewall rules. The minimum order value is 1 and the values must be unique.

**action**

Deny ▼

action can be Allow or Deny, default action is Allow.

protocolConfig

protocolConfig is a discriminated union of a protocol's specific configuration for TCP, UDP, SCTP, ICMP and ICMPv6. If not specified, packet matching will be based on the protocol value and protocol configuration, such as dstPort/type/code, will be ignored

protocol \*

ICMPv6 ▾

protocol can be ICMP, ICMPv6, TCP, SCTP or UDP.

icmp

icmp defines an ingress node firewall rule for ICMP protocol.

icmpv6

icmpv6 defines an ingress node firewall rule for ICMPv6 protocol.

icmpCode

114

icmpCode defines ICMP Code ID (RFC 792). if configured, this field matches against the ICMP/ICMPv6 header otherwise its ignored.

icmpType

128

icmpType defines ICMP Type Numbers (RFC 792). If configured, this field matches against the ICMP/ICMPv6 header otherwise its ignored.

[+ Add rules](#)

[+ Add ingress](#)

## interfaces

interfaces is a list of interfaces where the ingress firewall policy will be applied on.

Value •

eth0

[+ Add interfaces](#)

nodeSelector

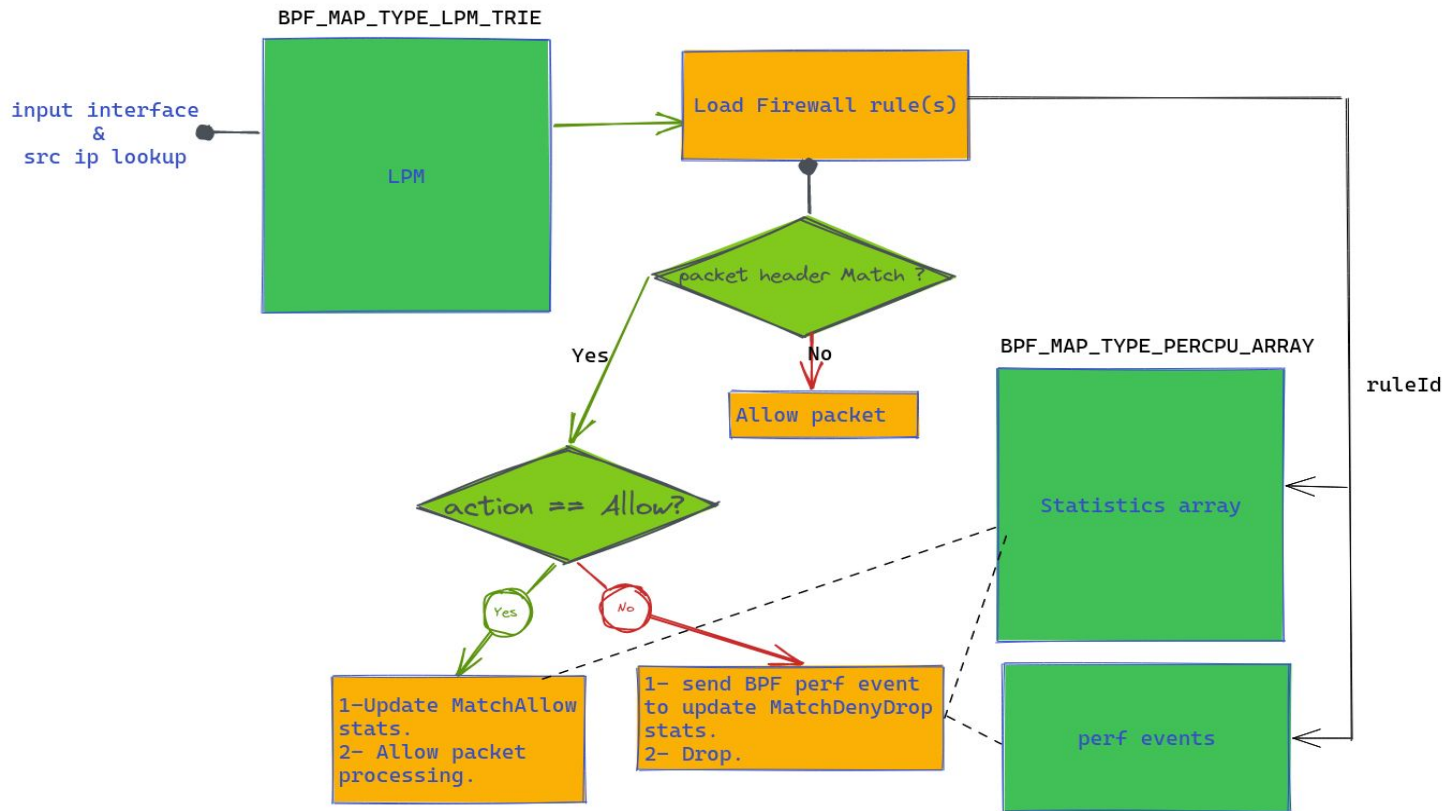
**nodeSelector** Selects node(s) where ingress firewall rules will be applied to.

matchExpressions

matchExpressions is a list of label selector requirements. The requirements are ANDed.

# Ingress node firewall XDP packet processing

CONFIDENTIAL Designator



# eBPF rules statistics sample output

```
oc exec -it -n openshift-ingress-node-firewall ingress-node-firewall-daemon-pqx56 -c daemon --
bash
# HELP ingressnodefirewall_node_packet_allow_bytes The number of bytes for packets which
results in an allow IP packet result
# TYPE ingressnodefirewall_node_packet_allow_bytes gauge
ingressnodefirewall_node_packet_allow_bytes 0
# HELP ingressnodefirewall_node_packet_allow_total The number of packets which results in an
allow IP packet result
# TYPE ingressnodefirewall_node_packet_allow_total gauge
ingressnodefirewall_node_packet_allow_total 0
# HELP ingressnodefirewall_node_packet_deny_bytes The number of bytes for packets which results
in an deny IP packet result
# TYPE ingressnodefirewall_node_packet_deny_bytes gauge
ingressnodefirewall_node_packet_deny_bytes 98
# HELP ingressnodefirewall_node_packet_deny_total The number of packets which results in a deny
IP packet result
# TYPE ingressnodefirewall_node_packet_deny_total gauge
ingressnodefirewall_node_packet_deny_total 1
```

## eBPF events log sample output

```
oc logs -n openshift-ingress-node-firewall ingress-node-firewall-daemon-pqx56 -c events
2022-08-25 14:50:41 +0000 UTC cnfdt19.lab.eng.tlv2.redhat.com ruleId 1 action Drop len 98 if eno1
2022-08-25 14:50:41 +0000 UTC cnfdt19.lab.eng.tlv2.redhat.com      ipv4 src addr 10.46.55.223 dst
addr 10.46.55.33
2022-08-25 14:50:41 +0000 UTC cnfdt19.lab.eng.tlv2.redhat.com      icmpv4 type 8 code 0
2022-08-25 15:01:34 +0000 UTC cnfdt19.lab.eng.tlv2.redhat.com ruleId 1 action Drop len 98 if eno1
2022-08-25 15:01:34 +0000 UTC cnfdt19.lab.eng.tlv2.redhat.com      ipv4 src addr 10.46.55.223 dst
addr 10.46.55.33
2022-08-25 15:01:34 +0000 UTC cnfdt19.lab.eng.tlv2.redhat.com      icmpv4 type 8 code 0
```

## Development Testing

- Component level unit-tests for all components.
- E2E testing coverage.

- oc adm must-gather – **gather\_ingress\_node\_firewall**
- From sos node's report it will collect eBPF bpftool outputs at */sos\_commands/ebpf*
- To enable XDP lookup matching debug we add debug config to IngressNodeFirewallConfig object and redeploy

```
cat <<EOF | oc apply -f -
apiVersion: ingressnodefirewall.openshift.io/v1alpha1
kind: IngressNodeFirewallConfig
metadata:
  name: ingressnodefirewallconfig
  namespace: openshift-ingress-node-firewall
spec:
  nodeSelector:
    node-role.kubernetes.io/worker: ""
  debug: true
EOF
```

This will create a hashmap table storing the Firewall rule lookup key and the value  
Use bpftool to dump this map and match the key with the packet's header

- Stateful Ingress node firewall to detect and block flood attacks :-
  - SYN Flood attacks
  - UDP Flood attacks
  - ICMP/ICMPv6 Flood attacks
- Large MTU support (Fixed in 4.14).
- Ability to chain multiple XDP programs to the same interface(s).
- Integrate with BPFd.



# DEMO

[Ingress node firewall demo on OCP cluster](#)